

## Lecture 2: Overfitting and Regularization

- Overfitting
- Cross-validation
- L2 and L1 regularization for linear estimators
- Bias-variance trade-off

## Recall: Steps to solving a supervised learning problem

1. Decide what the input-output pairs are.
2. Decide how to encode inputs and outputs.  
This defines the input space  $\mathcal{X}$ , and the output space  $\mathcal{Y}$ .
3. Choose a class of hypotheses/representations  $\mathcal{H}$ .
4. Choose an error function (cost function) to define the best hypothesis
5. Choose an algorithm for searching efficiently through the space of hypotheses.

## Recall: Linear hypothesis

- Suppose  $y$  was a linear function of  $\mathbf{x}$ :

$$h_{\mathbf{w}}(\mathbf{x}) = w_0 + w_1x_1(+\cdots)$$

- $w_i$  are called *parameters* or *weights*
- To simplify notation, we can add an attribute  $x_0 = 1$  to the other  $n$  attributes (also called *bias term* or *intercept term*):

$$h_{\mathbf{w}}(\mathbf{x}) = \sum_{i=0}^n w_i x_i = \mathbf{w}^T \mathbf{x}$$

where  $\mathbf{w}$  and  $\mathbf{x}$  are vectors of size  $n + 1$ .

## Recall: Least mean squares (LMS)


- Main idea: try to make  $h_{\mathbf{w}}(\mathbf{x})$  close to  $y$  on the examples in the training set
- We define a *sum-of-squares* error function

$$J(\mathbf{w}) = \frac{1}{2} \sum_{i=1}^m (h_{\mathbf{w}}(\mathbf{x}_i) - y_i)^2$$

(the  $1/2$  is just for convenience)

- We want to choose  $\mathbf{w}$  such as to minimize  $J(\mathbf{w})$

## Notation reminder: Gradient



- Multivariate generalization of the derivative.
- Points in the direction of the greatest increase of the function.

- Consider a function  $f(u_1, u_2, \dots, u_n) : \mathbb{R}^n \mapsto \mathbb{R}$  (e.g. an error function)
- The *partial derivative* w.r.t.  $u_i$  is denoted:

$$\frac{\partial}{\partial u_i} f(u_1, u_2, \dots, u_n) : \mathbb{R}^n \mapsto \mathbb{R}$$


The partial derivative is the derivative along the  $u_i$  axis, keeping all other variables fixed.

- The *gradient*  $\nabla f(u_1, u_2, \dots, u_n) : \mathbb{R}^n \mapsto \mathbb{R}^n$  is a function which outputs a vector containing the partial derivatives.

That is:

$$\nabla f = \left\langle \frac{\partial}{\partial u_1} f, \frac{\partial}{\partial u_2} f, \dots, \frac{\partial}{\partial u_n} f \right\rangle$$

# Properties of the gradient



- The inner product  $\langle \nabla f(\mathbf{x}), \mathbf{v} \rangle$  between the gradient of  $f$  at  $\mathbf{x}$  and any unit vector  $\mathbf{v} \in \mathbb{R}^n$  is the *directional derivative* of  $f$  in the direction of  $\mathbf{v}$  (i.e. the rate at which  $f$  changes at  $\mathbf{x}$  in the direction  $\mathbf{v}$ ).
  - $\nabla f(\mathbf{x})$  points towards the direction of greatest increase of  $f$  at  $\mathbf{x}$ .
  - Points such that  $\nabla f(\mathbf{x}) = \mathbf{0}$  are called **stationary points**.
  - The gradient  $\nabla f(\mathbf{x})$  is orthogonal to the contour line passing through  $\mathbf{x}$ .

## A bit of algebra

$$\begin{aligned}\frac{\partial}{\partial w_j} J(\mathbf{w}) &= \frac{\partial}{\partial w_j} \frac{1}{2} \sum_{i=1}^m (h_{\mathbf{w}}(\mathbf{x}_i) - y_i)^2 \\ &= \frac{1}{2} \cdot 2 \sum_{i=1}^m (h_{\mathbf{w}}(\mathbf{x}_i) - y_i) \frac{\partial}{\partial w_j} (h_{\mathbf{w}}(\mathbf{x}_i) - y_i) \\ &= \sum_{i=1}^m (h_{\mathbf{w}}(\mathbf{x}_i) - y_i) \frac{\partial}{\partial w_j} \left( \sum_{l=0}^n w_l x_{i,l} - y_i \right) \\ &= \sum_{i=1}^m (h_{\mathbf{w}}(\mathbf{x}_i) - y_i) x_{i,j}\end{aligned}$$

Setting all these partial derivatives to 0, we get a linear system with  $(n + 1)$  equations and  $(n + 1)$  unknowns.



## The solution

- Concise form of the error function:  $J(\mathbf{w}) = \frac{1}{2} \|\mathbf{X}\mathbf{w} - \mathbf{y}\|_2^2$ .
- Recalling some multivariate calculus:

$$\begin{aligned}\nabla_{\mathbf{w}} J &= \nabla_{\mathbf{w}} \frac{1}{2} (\mathbf{X}\mathbf{w} - \mathbf{y})^T (\mathbf{X}\mathbf{w} - \mathbf{y}) \\ &= \frac{1}{2} \nabla_{\mathbf{w}} (\mathbf{w}^T \mathbf{X}^T \mathbf{X} \mathbf{w} - \mathbf{y}^T \mathbf{X} \mathbf{w} - \mathbf{w}^T \mathbf{X}^T \mathbf{y} + \mathbf{y}^T \mathbf{y}) \\ &= \mathbf{X}^T \mathbf{X} \mathbf{w} - \mathbf{X}^T \mathbf{y} = \left[ \sum_{i=1}^m (h_{\mathbf{w}}(\mathbf{x}_i) - y_i) x_{i,j} \right]_{j=1 \dots n+1}\end{aligned}$$


- Setting gradient equal to zero:

$$\mathbf{X}^T \mathbf{X} \mathbf{w} - \mathbf{X}^T \mathbf{y} = 0 \quad \Rightarrow \quad \mathbf{X}^T \mathbf{X} \mathbf{w} = \mathbf{X}^T \mathbf{y} \quad \Rightarrow \quad \mathbf{w} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$$

- The inverse exists if the columns of  $\mathbf{X}$  are linearly independent.


## Example: Data and best linear hypothesis

$$y = 1.60x + 1.05$$



## Coming back to mean-squared error function...


- Good intuitive feel (small errors are ignored, large errors are penalized)
- Nice math (closed-form solution, unique global optimum)
- Geometric interpretation



## Linear regression summary

- The optimal solution (minimizing sum-squared-error) can be computed in polynomial time in the size of the data set.
- The solution is  $\mathbf{w} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$ , where  $\mathbf{X}$  is the data matrix augmented with a column of ones, and  $\mathbf{y}$  is the column vector of target outputs.
- A very rare case in which an analytical, exact solution is possible

## Predicting recurrence time based on tumor size



## Is linear regression enough?

- Linear regression is too simple for most realistic problems  
But it should be the first thing you try for real-valued outputs!
- Two possible solutions:
  1. Transform the data
    - Add cross-terms, higher-order terms
    - More generally, apply a transformation of the inputs from  $\mathcal{X}$  to some other space  $\mathcal{X}'$ , then do linear regression in the transformed space
  2. Use a different hypothesis class (e.g. non-linear functions)
- Today we focus on the first approach

## Polynomial fits

- Suppose we want to fit a higher-degree polynomial to the data (e.g.,  $y = w_0 + w_1x^1 + w_2x^2$ ).
- Suppose for now that there is a single input variable per training sample.
- How do we do it?

## Answer: Polynomial regression

- Given data:  $(x_1, y_1), (x_2, y_2), \dots, (x_m, y_m)$ .
- Suppose we want a degree- $d$  polynomial fit for the data

$$\mathbf{X} = [x_1, \dots, x_m]^\top, \quad \mathbf{y} = [y_1, \dots, y_m]^\top$$

- Let  $\mathbf{y}$  be as before and let

$$\Phi = \begin{bmatrix} 1 & x_1 & x_1^2 & \dots & x_1^d \\ 1 & x_2 & x_2^2 & \dots & x_2^d \\ \vdots & & \vdots & \ddots & \vdots \\ 1 & x_m & x_m^2 & \dots & x_m^d \end{bmatrix}$$

- Solve the linear regression  $\Phi \mathbf{w} \approx \mathbf{y}$ .



## Linear function approximation in general

- The best  $\mathbf{w}$  is considered the one which minimizes the sum-squared error over the training data:


$$\sum_{i=1}^m (y_i - h_{\mathbf{w}}(\mathbf{x}_i))^2$$

- We can find the best  $\mathbf{w}$  in closed form:

$$\mathbf{w} = (\Phi^T \Phi)^{-1} \Phi^T \mathbf{y}$$


- By linear models, we mean that the hypothesis function  $h_{\mathbf{w}}(\mathbf{x})$  is a linear function of the parameters  $\mathbf{w}$

## Order-2 fit




Is this a better fit to the data?

## Order-3 fit




Is this a better fit to the data?

## Order-4 fit




Is this a better fit to the data?

## Order-5 fit




Is this a better fit to the data?

## Order-6 fit




Is this a better fit to the data?

## Order-7 fit



Is this a better fit to the data?


## Order-8 fit



Is this a better fit to the data?



## Order-9 fit

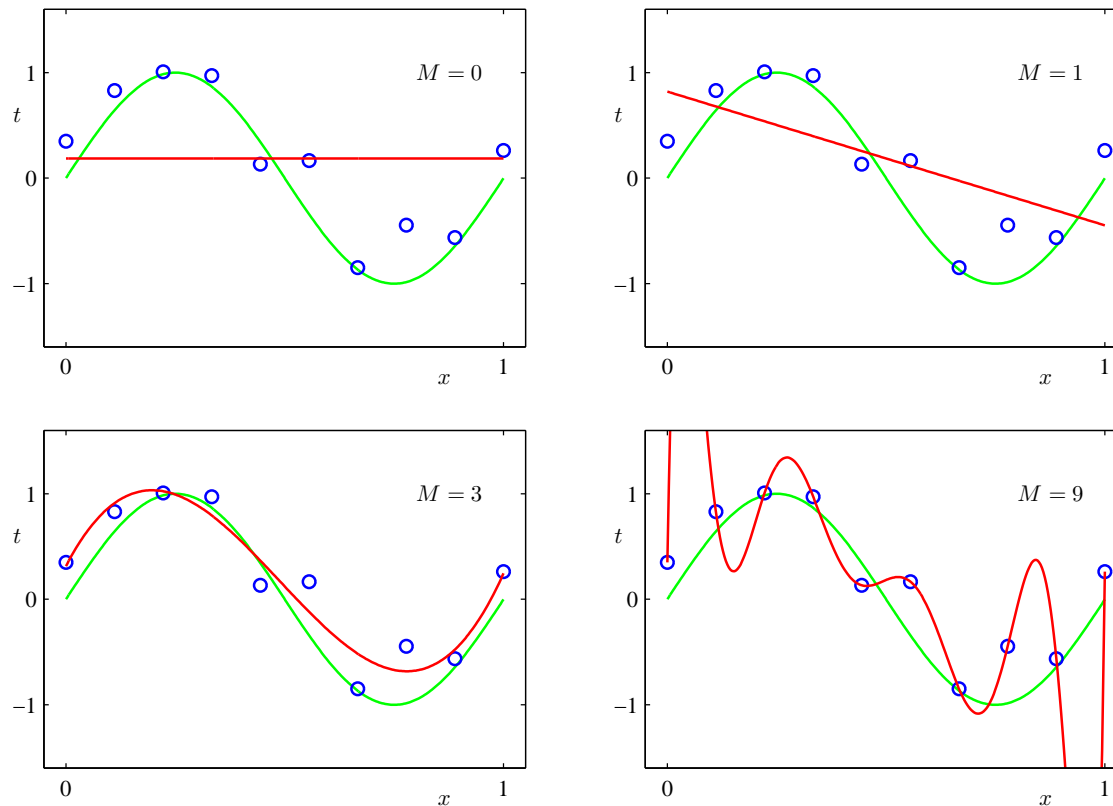


Is this a better fit to the data?

# Overfitting

- A general, HUGELY IMPORTANT problem for all machine learning algorithms
- We can find a hypothesis that predicts perfectly the training data but *does not generalize* well to new data
- E.g., a lookup table!
- We are seeing an instance here: if we have a lot of parameters, the hypothesis *memorizes* the data points, but is wild everywhere else.

## Another overfitting example

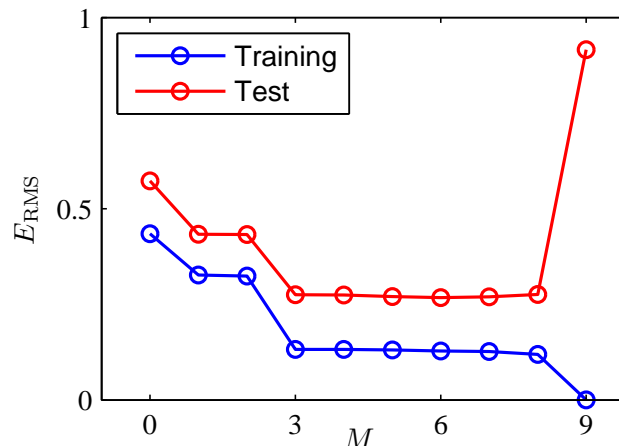


- The higher the degree of the polynomial  $M$ , the more degrees of freedom, and the more capacity to “overfit” the training data
- Typical overfitting means that error on the training data is very low, but error on new instances is high

## Overfitting more formally

- Assume that the data is drawn from some fixed, unknown probability distribution
- Every hypothesis has a *true* error  $J^*(h)$ , which is the expected error when data is drawn from the distribution.
- Because we do not have all the data, we measure the error on the training set  $J_D(h)$
- Suppose we compare hypotheses  $h_1$  and  $h_2$  on the training set, and  $J_D(h_1) < J_D(h_2)$
- If  $h_2$  is *truly* better, i.e.  $J^*(h_2) < J^*(h_1)$ , our algorithm is overfitting.
- We need theoretical and empirical methods to guard against it!

## Typical overfitting plot



- The training error decreases with the degree of the polynomial  $M$ , i.e. *the complexity of the hypothesis*
- The testing error, measured on independent data, decreases at first, then starts increasing
- Cross-validation helps us:
  - Find a good hypothesis class ( $M$  in our case), using a *validation set of data*
  - Report unbiased results, using a *test set*, untouched during either parameter training or validation

# Cross-validation

- A general procedure for estimating the true error of a predictor
- The available data is split into two subsets:
  - A *training and validation set* used only to find the right predictor
  - A *test set* used to report the prediction error of the algorithm
- These sets *must be disjoint!*
- The process is repeated several times, and the results are averaged to provide error estimates.

# Model selection with leave-one-out cross-validation

1. For each order of polynomial,  $d$ :
  - (a) Repeat the following procedure  $m$  times:
    - i. Leave out  *$i$ th instance* from the training set, to estimate the true prediction error; we will put it in a *validation set*
    - ii. Use all the other instances to find best parameter vector,  $\mathbf{w}_{d,i}$
    - iii. Measure the error in predicting the label on the instance left out, for the  $\mathbf{w}_{d,i}$  parameter vector; call this  $J_{d,i}$
    - iv. This is a *(mostly) unbiased estimate of the true prediction error*
  - (b) Compute the average of the estimated errors:  $J_d = \frac{1}{m} \sum_{i=1}^m J_{d,i}$
2. Choose the  $d$  with lowest average estimated error:  $d^* = \arg \min_d J(d)$

## Summary of leave-one-out cross-validation

- A very easy to implement algorithm
- Provides a great estimate of the true error of a predictor
- Computational cost scales with the number of instances (examples), so it can be prohibitive, especially if finding the best predictor is expensive
- Alternatives:
  - Leave- $k$ -out generalizes LOO, computationally very expensive.
  - $k$ -fold cross-validation: split the data set into  $k$  parts, then proceed as above.



# Regularization

- Remember the intuition: complicated hypotheses lead to overfitting
- Idea: change the error function to *penalize hypothesis complexity*:

$$J(\mathbf{w}) = J_D(\mathbf{w}) + \lambda J_{pen}(\mathbf{w})$$

This is called *regularization* in machine learning and *shrinkage* in statistics

- $\lambda$  is called *regularization coefficient* and controls how much we value fitting the data well, vs. a simple hypothesis

## $L_2$ regularization for linear models

- $L_2$  regularization (or *weight decay* in neural networks): add a squared penalty on the weights:

$$J_\lambda(\mathbf{w}) = \frac{1}{2}(\Phi\mathbf{w} - \mathbf{y})^\top(\Phi\mathbf{w} - \mathbf{y}) + \frac{\lambda}{2}\mathbf{w}^\top\mathbf{w}$$

- Why?
  - A simple hypothesis should not be too sensitive to small perturbation of the input
  - Math works out nicely
  - Resolve the issue of  $\Phi^\top\Phi$  not being invertible in linear regression, e.g. large  $n$  small  $m$  (original motivation)...

## $L_2$ regularization: closed form solution

$$J_\lambda(\mathbf{w}) = \frac{1}{2}(\Phi\mathbf{w} - \mathbf{y})^\top(\Phi\mathbf{w} - \mathbf{y}) + \frac{\lambda}{2}\mathbf{w}^\top\mathbf{w}$$

- By re-grouping terms, we get:

$$J_\lambda(\mathbf{w}) = \frac{1}{2}(\mathbf{w}^\top(\Phi^\top\Phi + \lambda\mathbf{I})\mathbf{w} - \mathbf{w}^\top\Phi^\top\mathbf{y} - \mathbf{y}^\top\Phi\mathbf{w} + \mathbf{y}^\top\mathbf{y})$$

- Optimal solution (obtained by solving  $\nabla J_\lambda(\mathbf{w}) = 0$ )

$$\mathbf{w} = (\Phi^\top\Phi + \lambda\mathbf{I})^{-1}\Phi^\top\mathbf{y}$$

(observe that  $\Phi^\top\Phi + \lambda\mathbf{I}$  is invertible)

## Side note: data centering


- For linear regression, we incorporated the bias term in  $\mathbf{w} \in \mathbb{R}^{n+1}$  by adding a 1 to each input vector  $\mathbf{x} \in \mathbb{R}^n$ . Why cannot we do the same for regression with  $L_2$  regularization?
  - Instead we center the data, i.e. remove means from inputs and outputs:
    - Let  $\bar{\mathbf{x}} = \frac{1}{m} \sum_{i=1}^m \mathbf{x}_i$  and  $\bar{y} = \frac{1}{m} \sum_{i=1}^m y_i$ .
    - Consider the new data set  $\{(\mathbf{x}_i - \bar{\mathbf{x}}, y_i - \bar{y})\}_{i=1}^m$ .
    - Solving the regularized regression problem on this new dataset is *equivalent* to solving the original problem
    - You can check that the bias term of the solution on this new data set is 0...
- ↪ exercise
- Also make sure that columns of  $\mathbf{X}$  are on the same scale (e.g. normalize)

## What $L_2$ regularization does

$$\arg \min_{\mathbf{w}} \frac{1}{2}(\Phi \mathbf{w} - \mathbf{y})^\top (\Phi \mathbf{w} - \mathbf{y}) + \frac{\lambda}{2} \mathbf{w}^\top \mathbf{w} = (\Phi^\top \Phi + \lambda \mathbf{I})^{-1} \Phi^\top \mathbf{y}$$

- If  $\lambda = 0$ , the solution is the same as in regular least-squares linear regression
- If  $\lambda \rightarrow \infty$ , the solution  $\mathbf{w} \rightarrow 0$
- Positive  $\lambda$  will cause the magnitude of the weights to be smaller than in the usual linear solution
- This is also known as *ridge regression*, and a special case of Tikhonov regularization

## Visualizing regularization (2 parameters)



$$\mathbf{w}^* = (\Phi^\top \Phi + \lambda \mathbf{I})^{-1} \Phi \mathbf{y}$$

## Pros and cons of $L_2$ regularization

- If  $\lambda$  is at a “good” value, regularization helps to avoid overfitting
- Choosing  $\lambda$  may be hard: cross-validation is often used
- If there are irrelevant features in the input (i.e. features that do not affect the output),  $L_2$  will give them small, but non-zero weights.
- Ideally, irrelevant input should have weights exactly equal to 0.

## $L_1$ Regularization for linear models

- Instead of requiring the  $L_2$  norm of the weight vector to be bounded, make the requirement on the  $L_1$  norm:

$$\min_{\mathbf{w}} J_D(\mathbf{w}) = \min_{\mathbf{w}} (\Phi \mathbf{w} - \mathbf{y})^\top (\Phi \mathbf{w} - \mathbf{y})$$

such that  $\sum_{i=1}^n |w_i| \leq \eta$

- This yields an algorithm called Lasso (Tibshirani, 1996)




## Solving $L_1$ regularization

- The optimization problem is a quadratic program
- There is one constraint for each possible sign of the weights ( $2^n$  constraints for  $n$  weights)
- For example, with two weights:

$$\begin{aligned} \min_{w_1, w_2} \quad & \sum_{j=1}^m (y_j - w_1 x_1 - w_2 x_2)^2 \\ \text{such that } & w_1 + w_2 \leq \eta \\ & w_1 - w_2 \leq \eta \\ & -w_1 + w_2 \leq \eta \\ & -w_1 - w_2 \leq \eta \end{aligned}$$

- Solving this program directly can be done for problems with a small number of inputs

## Visualizing $L_1$ regularization





- If  $\lambda$  is big enough, the circle is very likely to intersect the diamond at one of the corners
- This makes  $L_1$  regularization much more likely to make some weights *exactly* 0

## Pros and cons of $L_1$ regularization

- If there are irrelevant input features, Lasso is likely to make their weights 0, while  $L_2$  is likely to just make all weights small
- Lasso is biased towards providing *sparse solutions* in general
- Lasso optimization is computationally more expensive than  $L_2$
- More efficient solution methods have to be used for large numbers of inputs (e.g. least-angle regression, 2003).
- $L_1$  methods of various types are very popular
- One can combine  $L_1$  and  $L_2$  regularization (elastic-net)

## Example of L1 vs L2 effect



- Note the sparsity in the coefficients induced by  $L_1$

## The anatomy of the error of an estimator

- Suppose we have a dataset  $\{(\mathbf{x}_i, y_i)\}_{i=1}^m$  where  $y = f(\mathbf{x}) + \epsilon$  and  $\epsilon$  is Gaussian noise with zero mean and standard deviation  $\sigma^2$
- We fit a linear hypothesis  $h(\mathbf{x}) = \mathbf{w}^\top \mathbf{x}$ , such as to minimize sum-squared error over the training data:

$$\sum_{i=1}^m (y_i - h(\mathbf{x}_i))^2$$

- Because of the hypothesis class that we chose (hypotheses linear in the parameters) for some target functions  $f$  we will have a *systematic prediction error*
- Even if  $f$  were truly from the hypothesis class we picked, depending on the data set we have, the parameters  $\mathbf{w}$  that we find may be different; this *variability* due to the specific data set on hand is a different source of error

## Bias-variance analysis

- Given a new data point  $\mathbf{x}$ , what is the *expected prediction error*?
- Assume that the data points are drawn *independently and identically distributed (i.i.d.)* from a unique underlying probability distribution  $P(\mathbf{x}, y) = P(\mathbf{x})P(y|\mathbf{x})$
- The goal of the analysis is to compute, for an arbitrary given point  $\mathbf{x}$ ,

$$E_P [(y - h(\mathbf{x}))^2 | \mathbf{x}]$$

where the expectation is over all training sets of a given size drawn according to  $P(\cdot, \cdot)$  and over  $y$  drawn according to  $P(\cdot | \mathbf{x})$ .

- For a given hypothesis class, we can also compute the *true error*, which is the expected error over the input distribution:

$$\sum_{\mathbf{x}} E_P [(y - h(\mathbf{x}))^2 | \mathbf{x}] P(\mathbf{x})$$

(if  $\mathbf{x}$  continuous, sum becomes integral with appropriate conditions).

- We will decompose this expectation into three components

## Recall: Statistics 101

- Let  $X$  be a random variable with possible values  $x_i, i = 1 \dots n$  and with probability distribution  $P(X)$
- The *expected value* or *mean* of  $X$  is:

$$E[X] = \sum_{i=1}^n x_i P(x_i)$$

- If  $X$  is continuous, roughly speaking, the sum is replaced by an integral, and the distribution by a density function
- The *variance* of  $X$  is:

$$\begin{aligned} Var[X] &= E[(X - E(X))^2] \\ &= E[X^2] - (E[X])^2 \end{aligned}$$

## The variance lemma

$$\begin{aligned} \text{Var}[X] &= E[(X - E[X])^2] \\ &= \sum_{i=1}^n (x_i - E[X])^2 P(x_i) \\ &= \sum_{i=1}^n (x_i^2 - 2x_i E[X] + (E[X])^2) P(x_i) \\ &= \sum_{i=1}^n x_i^2 P(x_i) - 2E[X] \sum_{i=1}^n x_i P(x_i) + (E[X])^2 \sum_{i=1}^n P(x_i) \\ &= E[X^2] - 2E[X]E[X] + (E[X])^2 \cdot 1 \\ &= E[X^2] - (E[X])^2 \end{aligned}$$

We will use the form:

$$E[X^2] = (E[X])^2 + \text{Var}[X]$$



## Bias-variance decomposition

- Simple algebra:

$$\begin{aligned} E_P [(y - h(\mathbf{x}))^2 | \mathbf{x}] &= E_P [(h(\mathbf{x}))^2 - 2yh(\mathbf{x}) + y^2 | \mathbf{x}] \\ &= E_P [(h(\mathbf{x}))^2 | \mathbf{x}] + E_P [y^2 | \mathbf{x}] - 2E_P [y | \mathbf{x}] E_P [h(\mathbf{x}) | \mathbf{x}] \end{aligned}$$

- Let  $\bar{h}(\mathbf{x}) = E_P[h(\mathbf{x}) | \mathbf{x}]$  denote the *mean prediction* of the hypothesis at  $\mathbf{x}$ , when  $h$  is trained with data drawn from  $P$
- For the first term, using the variance lemma, we have:

$$E_P[(h(\mathbf{x}))^2 | \mathbf{x}] = E_P[(h(\mathbf{x}) - \bar{h}(\mathbf{x}))^2 | \mathbf{x}] + (\bar{h}(\mathbf{x}))^2$$

- Note that  $E_P[y | \mathbf{x}] = E_P[f(\mathbf{x}) + \epsilon | \mathbf{x}] = f(\mathbf{x})$  (because of linearity of expectation and the assumption on  $\epsilon \sim \mathcal{N}(0, \sigma)$ )
- For the second term, using the variance lemma, we have:

$$E[y^2 | \mathbf{x}] = E[(y - f(\mathbf{x}))^2 | \mathbf{x}] + (f(\mathbf{x}))^2$$


## Bias-variance decomposition (2)

- Putting everything together, we have:

$$\begin{aligned} E_P [(y - h(\mathbf{x}))^2 | \mathbf{x}] &= E_P [(h(\mathbf{x}) - \bar{h}(\mathbf{x}))^2 | \mathbf{x}] + (\bar{h}(\mathbf{x}))^2 - 2f(\mathbf{x})\bar{h}(\mathbf{x}) \\ &\quad + E_P [(y - f(\mathbf{x}))^2 | \mathbf{x}] + (f(\mathbf{x}))^2 \\ &= E_P [(h(\mathbf{x}) - \bar{h}(\mathbf{x}))^2 | \mathbf{x}] + (f(\mathbf{x}) - \bar{h}(\mathbf{x}))^2 + E[(y - f(\mathbf{x}))^2 | \mathbf{x}] \end{aligned}$$

- The first term,  $E_P [(h(\mathbf{x}) - \bar{h}(\mathbf{x}))^2 | \mathbf{x}]$ , is the *variance* of the hypothesis  $h$  at  $\mathbf{x}$ , when trained with finite data sets sampled randomly from  $P$
- The second term,  $(f(\mathbf{x}) - \bar{h}(\mathbf{x}))^2$ , is the *squared bias* (or systematic error) which is associated with the class of hypotheses we are considering
- The last term,  $E[(y - f(\mathbf{x}))^2 | \mathbf{x}]$  is the *noise*, which is due to the problem at hand, and cannot be avoided

# Error decomposition



- The bias-variance sum approximates well the test error over a set of 1000 points
- x-axis measures the hypothesis complexity (decreasing left-to-right)
- Simple hypotheses usually have high bias (bias will be high at many points, so it will likely be high for many possible input distributions)
- Complex hypotheses have high variance: the hypothesis is very dependent on the data set on which it was trained.

# Bias-variance trade-off




image credit: Scott Fortman-roe (<http://scott.fortmann-roe.com/docs/BiasVariance.html>)

## Bias-variance trade-off

- Typically, bias comes from not having good hypotheses in the considered class
- Variance results from the hypothesis class containing “too many” hypotheses
- Hence, we are faced with a *trade-off*: choose a more expressive class of hypotheses, which will generate higher variance, or a less expressive class, which will generate higher bias
- Making the trade-off has to depend on the amount of data available to fit the parameters (data usually mitigates the variance problem)

## More on overfitting

- Overfitting depends on the amount of data, relative to the complexity of the hypothesis
- With more data, we can explore more complex hypotheses spaces, and still find a good solution

